

Deep Learning

Exercise 6: Convolutional Networks

Room: **BIN-1-B.01**

Instructor: Manuel Günther

Email: guenther@ifi.uzh.ch

Office: AND 2.54

Friday, April 1, 2022

Outline

- 1 PyTorch
- 2 MNIST Training with PyTorch

Outline

- 1 PyTorch
 - Types of Layers
 - Datasets and Batches
 - The Training Loop
 - Running on the GPU

Types of Layers

Learnable Layers

- Fully connected layer
`torch.nn.Linear`
 - `in_features = D`
 - `out_features = K`
 - `bias = True`
- Convolutional layer
`torch.nn.Conv2d`
 - `in_channels = C`
 - `out_channels = Q`
 - `kernel_size = U` or (U, V)
 - `padding, stride, bias`

Non-Learnable Layers

- Activation functions:
`torch.nn.Sigmoid`,
`torch.nn.Tanh`,
`torch.nn.Softmax` ...
- Pooling:
`torch.nn.MaxPool2d`,
`torch.nn.AvgPool2d`
 - `kernel_size, stride`
- Input flattening:
`torch.nn.Flatten`

Datasets and Batches

Datasets

- Many available in `torchvision.datasets`
 - `torchvision.datasets.MNIST`
 - `torchvision.datasets.ImageNet`
- Common interface:
 - `root`: Directory of raw data
 - `train`: set to `False` for test set
 - `download`: downloads data if required
 - `transform`: preprocessing
- Return `PIL` images

Transforms

- Prepares the input
 - Depends on original data
- Implemented in `torchvision.transforms`
 - `Resize((D, E))`: height, width
 - `Normalize(mean, std)`
 - `Lambda(callable)`: generic
 - `ToTensor`: `PIL` → `tensor`
- Combining several transforms:
 - `Compose((trans1,trans2))`
- Additionally: target transforms

Datasets and Batches

Data Loader

- Prepares batches of data
 - For both input and target
- `torch.utils.data.DataLoader`
 - `dataset`: see above
 - `batch_size` = B
 - `shuffle` after each epoch
 - `num_workers`: parallel execution on the CPU (might slow down processing)

Example MNIST Dataset

```
# obtain datasets
transform = torchvision.transforms.ToTensor()
train_set = torchvision.datasets.MNIST(
    root="/temp/MNIST",
    train=True, download=True,
    transform=transform
)
test_set = torchvision.datasets.MNIST(
    root="/temp/MNIST",
    train=False, download=True,
    transform=transform
)

# loaders
train_loader = torch.utils.data.DataLoader(
    train_set, shuffle=True, batch_size=64
)
test_loader = torch.utils.data.DataLoader(
    test_set, shuffle=False, batch_size=100
)
```

The Training Loop

Example Training Loop

```
# training loop
for epoch in range(epochs):
    for x,t in train_loader:
        # DO NOT FORGET:
        optimizer.zero_grad()
        z = network(x)
        J = loss(z, t)
        J.backward()
        optimizer.step()

    # compute train accuracy
    ...
```

Example Test Loop

```
# testing loop
with torch.no_grad():
    correct = 0
    for x,t in test_loader:
        z = network(x)
        # optional: compute test loss
        J = loss(z, t)

        # compute test accuracy
        correct += torch.sum(
            torch.argmax(z, dim=1) == t
        ).item()
    acc = correct/ len(test_set)
```

Running on the GPU

Preparation

- Test CUDA availability:
`torch.cuda.is_available()`
- Select device "cpu" or "cuda":
`device = torch.device("cuda")`
- Move everything to the device:
`network.to(device)`
`x.to(device)`
`t.to(device)`

Speed Warning

Can be slow on CPU (1 min per epoch)

Example Training Loop

```
# instantiate everything
device = torch.device("cuda")
network = Network().to(device)
loss = torch.nn.CrossEntropyLoss()
optimizer = torch.optim.SGD(
    params=network.parameters(),
    lr=0.01, momentum=0.9
)

# training epoch
for epoch in range(epochs):
    for x,t in train_loader:
        optimizer.zero_grad()
        z = network(x.to(device))
        J = loss(z, t.to(device))
        J.backward()
        optimizer.step()
```


Running on the GPU

Enabling GPU Processing on Google Colaboratory

- Google Colaboratory allows usage of Google GPU servers
- Open a notebook on Google Colaboratory
- Select Runtime → Change Runtime Type
- Select Acceleration GPU (does not work for TPU)
- Check that GPU is enabled:
 - `torch.cuda.is_available()` should return `True`

Outline

- 2 MNIST Training with PyTorch
 - Dataset and Data Loader
 - Networks
 - Network Training

MNIST Training with PyTorch

Goal of Exercise

- Compare fully-connected and convolutional networks
- Get to know PyTorch `Dataset` and `DataLoader` classes
- Familiarize with PyTorch training procedure with batches

Task 1: Dataset

- Implement a function to return two datasets with a given `transform`
- Instantiate MNIST dataset from `torchvision.datasets`
 - Set dataset to automatically download
- Provide both training and test set

Datasets and Data Loaders

Test 1: Data Types

- Instantiate datasets with `transform=None`
- Check that the dataset returns `PIL.Image.Image`'s

Task 2: Data Loaders

- Load datasets with `torchvision.transforms.ToTensor` transform
- Create two data loaders, one for each dataset
 - Choose training batch size to be $B = 64$ (B for test on your choice)
 - Decide which of the two data loaders require shuffling

Test 2: Batches

- Check data types and ranges of batches: input and target
- Check that batch size equals B for all except the last batch

Networks

Compared Networks

- Fully-connected and convolutional network
→ Same number of layers with weights

Fully-connected Network

- 1 Flatten layer to convert 28×28 input into $28 * 28$ vector
- 2 $K \times D$ fully-connected layer
- 3 tanh activation
- 4 $K \times K$ fully-connected layer
- 5 tanh activation
- 6 $K \times O$ fully-connected layer

Convolutional Network

- 1 $Q_1 \times 1 \times 5 \times 5$ convolutional layer, stride 1, padding 2
- 2 Maximum pooling, 2×2 kernel, stride 2
- 3 tanh activation
- 4 $Q_2 \times Q_1 \times 5 \times 5$ convolutional layer, stride 1, padding 2
- 5 Maximum pooling, 2×2 kernel, stride 2
- 6 tanh activation
- 7 Flatten layer
- 8 $O \times ?$ fully-connected layer

Networks

Task 3: Fully-Connected Network

- Implement a function to return the fully-connected network
 - `torch.nn.Sequential` can still be used

Task 4: Convolutions Output (theoretical question)

- Analytically compute input size of fully-connected layer in convolutional network
 - Consider kernel sizes, strides, paddings and poolings

Task 5: Convolutional Network

- Implement a function to return the convolutional network
 - Consider result of task 4
 - `torch.nn.Sequential` can be used, too

Network Training

Task 6: Training and Validation Loop

- `training` function taking `network`, `epochs`, `eta`
- Instantiate categorical cross-entropy `loss`
- Instantiate stochastic gradient decent `optimizer`
- For `epochs` iterate:
 - 1 Train on all batches of the training set
 - 2 Compute loss and accuracy on test set
 - Store both in separate lists
- Return both lists of losses and accuracies

Network Training

Task 7: Fully-connected Training

- Instantiate fully-connected network with $K = 100$ and $O = 10$
- Train fully-connected network for 10 epochs with $\eta = 0.01$
 - Store lists of losses and accuracies

Task 8: Convolutional Training



- Create convolutional network with $Q_1 = 32$, $Q_2 = 64$ and $O = 10$
- Train convolutional network for 10 epochs with $\eta = 0.01$
 - Store lists of losses and accuracies

Task 9: Plotting

- Plot losses and accuracies in two separate plots

Network Training

Exemplary Losses and Accuracies (for 100 Epochs)



Task 10: Learnable Parameters

- Compute number of learnable parameters for both networks
→ Both analytically and via PyTorch