

Deep Learning

Exercise 9: Convolutional Auto-Encoder

Room: **BIN-1-B.01**

Instructor: Manuel Günther

Email: guenther@ifi.uzh.ch

Office: AND 2.54

Friday, May 6, 2022

Outline

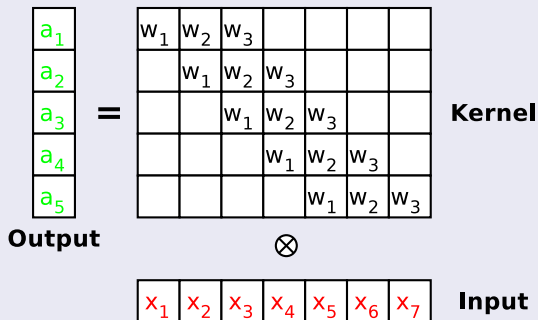
- 1 PyTorch
- 2 Convolutional Auto-Encoder

Outline

- 1 PyTorch
 - Fractionally-Strided Convolution
 - Network Implementation

Fractionally-Strided Convolution

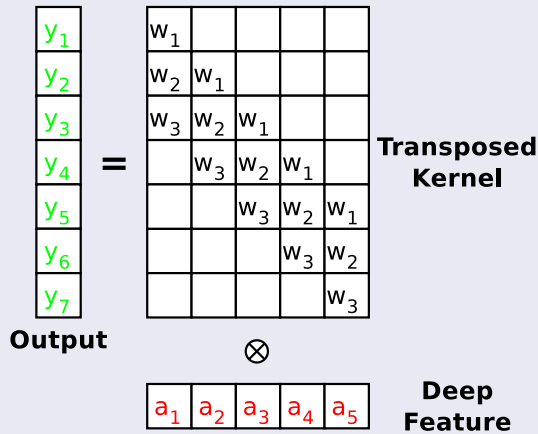
Convolution by Multiplication



Implementation

- Implemented: `torch.nn.Conv2d`

Transposed Convolution



Fractionally-Strided Convolution

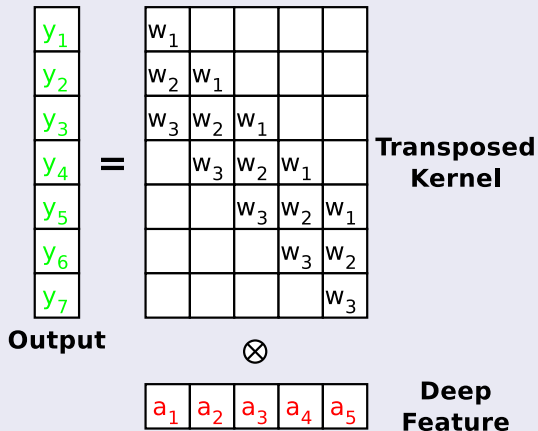
Implementation

- **Inverts** `Conv2d`
- `torch.nn.ConvTranspose2d`
- Parameters
 - `stride`: fractional stride
 - `padding`: same as `Conv2d`
 - `output_padding`

Output Padding

- For `stride=2`, output dimensions always odd
- Add output padding of 1

Transposed Convolution



Fractionally-Strided Convolution

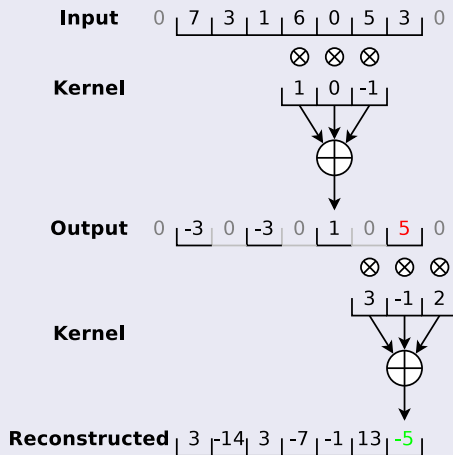
Implementation

- **Inverts** `Conv2d`
- `torch.nn.ConvTranspose2d`
- Parameters
 - `stride`: fractional stride
 - `padding`: same as `Conv2d`
 - `output_padding`

Output Padding

- For `stride=2`, output dimensions always odd
- Add output padding of 1

Fraction. Strided Conv.



Network Implementation

Reminder: Modules in PyTorch

- A module can be
 - A separate layer (e.g. Linear, ReLU, ...)
 - A block of layers (e.g. ResNet Block)
 - A complete network (e.g. LeNet, ResNet)
- ⇒ A network is a module

Building Cascaded Modules

- Submodules can be networks
- Auto-encoder: two networks
 - Encoder network $\vec{\varphi} = \mathcal{E}(\mathcal{X})$
 - Decoder network $\mathcal{Y} = \mathcal{D}(\vec{\varphi})$
- Creating auto-encoder:
 - Implement encoder module
 - Implement decoder module
 - Instantiate both in `__init__`
 - Cascade both in `forward`
 - Access through variable

Network Implementation

Encoder Network

```
class Encoder (torch.nn.Module):  
    def __init__(self, ...):  
        self.conv1 = torch.nn.Conv2d(...)
```

Decoder Network

```
class Decoder (torch.nn.Module):  
    def __init__(self, ...):  
        self.deconv1 = torch.nn.ConvTranspose2d(...)
```

Auto-Encoder Network

```
class AutoEncoder (torch.nn.Module):  
    def __init__(self, ...):  
        self.encoder = Encoder(...)  
        self.decoder = Decoder(...)
```


Network Implementation

Encoder Network

- 1 2D convolution: Q_1 channels, 5×5 kernel, stride 2, padding 2
- 2 activation function
- 3 2D convolution: Q_2 channels, 5×5 kernel, stride 2, padding 2
- 4 flatten layer
- 5 activation function
- 6 fully-connected layer: K outputs

Decoder Network

- 1 fully-connected layer: K inputs
- 2 activation function
- 3 reshaping
- 4 2D fractionally strided-convolution: Q_2 channels, 5×5 kernel, stride 2, padding 2
- 5 activation function
- 6 2D fractionally strided-convolution: Q_1 channels, 5×5 kernel, stride 2, padding 2

Decoder goes backward through encoder layers

Outline

- 2 Convolutional Auto-Encoder
 - Dataset
 - Auto-Encoder
 - Training and Evaluation

Dataset

Task 1: Dataset

- Rely on MNIST dataset, ignore labels
 - Except for evaluation Task 7 where we need labels
- Create training and validation/test sets
 - Simple `ToTensor` transform
- `DataLoaders`: $B = 32$ for training, $B = 100$ for validation
- Nothing special.

Auto-Encoder

Task 2: Encoder Network

- Implement convolutional `Encoder` network
- Provide constructor `__init__(self, ...)`
- Provide `forward(self, x)` implementation

Task 3: Decoder Network

- Implement fractionally strided-convolutional `Decoder` network
- Provide constructor `__init__(self, ...)`
- Provide `forward(self, x)` implementation
- Output values should be in range $[0, 1]$

Auto-Encoder

Task 4: Joint Auto-Encoder Network

- Implement `AutoEncoder`
- Combine `Encoder` and `Decoder`
- Implement `__init__(self, ...)` and `forward(self, x)`

Test 1: Output Sizes

- Instantiate auto-encoder network: $Q_1 = 32, Q_2 = 32, K = 10$
- Define input \mathbf{X} for network
- Compute output of network and check its size and value range
- Check `network.encoder` and `network.decoder` separately

Training and Evaluation

Task 5: Training and Validation Loop

- Use \mathcal{J}^{L_2} loss function `torch.nn.MSELoss`
 - Compare output \mathbf{Y} with input \mathbf{X}
- Use Adam optimizer with learning rate $\eta = 0.001$ (or lower)
- Train for 10 epochs
 - Compute training set loss during training
 - Compute and report validation set loss after epoch

When losses stagnate after epoch 2, reduce η !

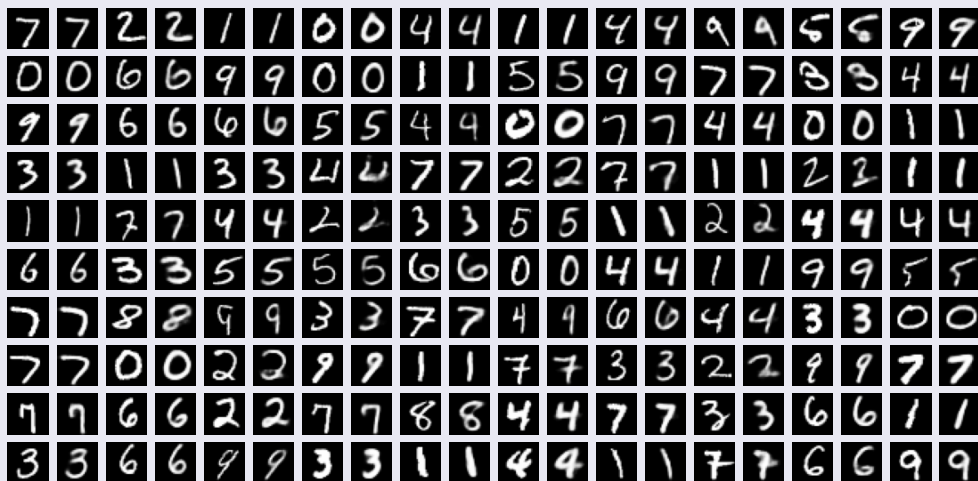
Training and Evaluation

Task 6: Reconstruction Result

- Take batch of images from test set
- Run batch through auto-encoder
- Plot input and output next to each other
 - 10 rows with 10 pairs of images

Training and Evaluation

Pairs of Inputs and Outputs of Auto-Encoder



Training and Evaluation

Task 7: Mean Vector per Class

- Use encoder to extract deep feature for validation samples
- Split them by class label
- Compute average feature $\vec{\mu}_o$ for each class o

Task 8: Decode Mixtures of Classes

- Iterate over all possible pairs (o_1, o_2) of classes
- Compute mixed deep feature $\frac{\vec{\mu}_{o_1} + \vec{\mu}_{o_2}}{2}$ of the two classes
→ If $o_1 = o_2 = o$, we get the class average $\vec{\mu}_o$
- Use decoder to generate output for deep feature
- Plot in a grid with 10 rows and 10 columns

Training and Evaluation

Reconstructed Mix of Classes

